

(iii) *Regression loss*  $\mathcal{L}_{\text{regress}}$ . For each pair  $(i, j)$  we have a ground-truth indirect score  $\text{trueindirectscore}_{ij} \in [0, 1]$ . Let

$$e_{ij} = I_{\text{learned}}(i, j) - \text{trueindirectscore}_{ij}.$$

We use the Huber loss with parameter  $\delta > 0$ :

$$\mathcal{L}_{\text{regress}}(i, j) = \begin{cases} \frac{1}{2}e_{ij}^2, & |e_{ij}| \leq \delta, \\ \delta|e_{ij}| - \frac{1}{2}\delta^2, & |e_{ij}| > \delta. \end{cases}$$

Averaging over all training pairs yields  $\mathcal{L}_{\text{regress}}$ .

*Output.* A learned path-aggregation model  $LPA_{\text{model}}^{(0)}$  consisting of a path encoder and an attention-based aggregator capable of mapping sets of paths to calibrated indirect scores.

### (c) The EM-Refinement Loop

**Goal.** The objective of this step is to refine the pre-trained (v0.1) models from Step 4.2 on the real, unlabeled document corpus by using their own predictions as a self-supervised signal. The refinement is formulated as an Expectation–Maximization (EM) style loop with a Mean-Teacher architecture to ensure stability.

We maintain two sets of parameters for each learnable component (e.g., fusion model, path aggregator):

- *Student parameters*  $\theta_S$ : actively updated by gradient-based optimization.
- *Teacher parameters*  $\theta_T$ : updated as an exponential moving average (EMA) of  $\theta_S$  and used to generate pseudo-labels.

For concreteness, let  $\theta_{S,\text{fusion}}$  and  $\theta_{T,\text{fusion}}$  denote the parameters of the student and teacher fusion models, and analogously for the LPA model.

We perform  $R$  refinement rounds:

$$\text{for } r = 1, \dots, R.$$

#### A. E-Step (Expectation / Pseudo-Label Estimation).

*Goal.* For a fixed round  $r$ , the E-step uses the current teacher models to compute the best available estimate of the causal answer key  $C_{\text{prior}}^{(r)}$  on the real corpus.

*Procedure.*

- We run the complete CoCaD pipeline (Phase 3: Steps 3.1–3.3) on the real document corpus, using the teacher parameter sets from the previous round,  $\theta_T^{(r-1)}$ , for all learnable modules (fusion model, LPA model, CPC model, retriever, etc.). This yields, for each candidate pair  $(i, j) \in E_{\text{prior}}$ , a final score and associated statistics:

$$(i, j) \mapsto (\text{finalscore}_{ij}^{(r)}, p(i, j), v_{ij}, H_p(i, j), \dots),$$

where  $p(i, j)$  is the empirical p-value from Step 3.3 and  $v_{ij}$  is the feature vector used by the fusion model.

- In round  $r = 1$ , the conditional LLM component in Step 3.1 is executed explicitly to obtain scores  $\mu_{\text{LLMcond}}(v_{ij})$  for all  $(i, j)$ , and we store the pairs

$$\mathcal{D}_{\text{distill}} = \{(v_{ij}, \mu_{\text{LLMcond}}(v_{ij}))\}.$$

- iii. In the M-step of round 1 (see below), we train a student surrogate model  $f_{\text{student}}$  to approximate the LLM scores. In subsequent rounds  $r > 1$ , the expensive LLM calls in Step 3.1 are deterministically replaced by  $f_{\text{student}}(v_{ij})$ , so that the E-step remains computationally feasible while preserving the interface of the pipeline.

The output of the E-step at round  $r$  is the sparse answer key

$$C_{\text{prior}}^{(r)} = \{(i, j, \text{finalscore}_{ij}^{(r)}, p(i, j), \dots)\}.$$

## B. M-Step (Maximization / Student Update).

*Goal.* For fixed pseudo-labels  $C_{\text{prior}}^{(r)}$ , the M-step updates the student parameters  $\theta_S^{(r)}$  by minimizing a combination of pseudo-label losses and consistency regularization on real data.

*Soft, Confidence-Weighted Pseudo-Labels.* For each pair  $(i, j)$  in  $C_{\text{prior}}^{(r)}$  we define:

$$\begin{aligned} y_{ij}^{\text{soft}} &= \text{finalscore}_{ij}^{(r)} \in [0, 1], \\ w_{ij} &= 1 - p(i, j), \end{aligned}$$

where  $w_{ij}$  acts as a confidence weight: highly significant links (small  $p(i, j)$ ) receive weights close to 1, whereas non-significant links receive weights close to 0.

We also define a hard binary label for pruning decisions,

$$y_{ij}^{\text{bin}} = \begin{cases} 1, & \text{if } \text{finalscore}_{ij}^{(r)} > 0, \\ 0, & \text{otherwise.} \end{cases}$$

### B.1. Fusion Model Update.

The student fusion model  $f_{\text{fusion}, S}^{(r)}$  with parameters  $\theta_{S, \text{fusion}}^{(r)}$  maps feature vectors  $v_{ij}$  to scores in  $(0, 1)$ . We update  $\theta_{S, \text{fusion}}^{(r)}$  by minimizing a confidence-weighted binary cross-entropy loss:

$$\mathcal{L}_{\text{fusion-real}} = \sum_{(i, j) \in C_{\text{prior}}^{(r)}} w_{ij} \cdot \text{BCE}(f_{\text{fusion}, S}^{(r)}(v_{ij}), y_{ij}^{\text{soft}}),$$

where  $\text{BCE}(p, y) = -(y \log p + (1 - y) \log(1 - p))$ .

Gradients of  $\mathcal{L}_{\text{fusion-real}}$  with respect to  $\theta_{S, \text{fusion}}^{(r)}$  are computed and applied using standard stochastic gradient descent or Adam.

### B.2. LPA Model Update.

The student LPA model  $LPA_S^{(r)}$ , with parameters  $\theta_{S, \text{LPA}}^{(r)}$ , outputs for each pair  $(i, j)$  an aggregated indirect score  $I_{\text{learned}, S}^{(r)}(i, j) \in (0, 1)$ . In the refinement loop we treat this as a classifier for pruning and update it using the binary pseudo-labels:

$$\mathcal{L}_{\text{LPA-real}} = \sum_{(i, j) \in C_{\text{prior}}^{(r)}} w_{ij} \cdot \text{BCE}(I_{\text{learned}, S}^{(r)}(i, j), y_{ij}^{\text{bin}}).$$

The parameters  $\theta_{S, \text{LPA}}^{(r)}$  are updated by backpropagation on this loss.

### B.3. Distillation of the Conditional LLM (First Round Only).

In round  $r = 1$ , we additionally train a student surrogate  $f_{\text{student}}$  on the distillation set  $\mathcal{D}_{\text{distill}}$ :

$$\mathcal{L}_{\text{distill}} = \sum_{(v_{ij}, t_{ij}) \in \mathcal{D}_{\text{distill}}} \|f_{\text{student}}(v_{ij}) - t_{ij}\|_2^2,$$

where  $t_{ij} = \mu_{\text{LLMcond}}(v_{ij})$  is the original LLM score. The weights of  $f_{\text{student}}$  are updated via gradient descent on  $\mathcal{L}_{\text{distill}}$ , and the trained  $f_{\text{student}}$  is then used in place of the LLM in all subsequent rounds.

### B.4. Consistency Regularization.

To improve robustness to perturbations and prevent overfitting to noisy pseudo-labels, we introduce a consistency loss for each student model  $f_{\text{student},S}^{(r)}$  (this notation can refer generically to the fusion model, LPA model, or other learnable components). For a given  $(i, j)$ , we construct two stochastic augmentations of the input,  $v_{ij}^{\text{aug1}}$  and  $v_{ij}^{\text{aug2}}$  (e.g., via feature dropout, small noise in embeddings) and enforce consistency between their predictions:

$$\mathcal{L}_{\text{consistency}} = \sum_{(i,j)} \|f_{\text{student},S}^{(r)}(v_{ij}^{\text{aug1}}) - f_{\text{student},S}^{(r)}(v_{ij}^{\text{aug2}})\|_2^2.$$

### B.5. Total M-Step Objective.

For each student model we define the total M-step loss as

$$\mathcal{L}_M = \mathcal{L}_{\text{pseudo-label}} + \lambda_{\text{consist}} \cdot \mathcal{L}_{\text{consistency}},$$

where  $\mathcal{L}_{\text{pseudo-label}}$  denotes the corresponding pseudo-label loss (e.g.,  $\mathcal{L}_{\text{fusion-real}}$  for the fusion model,  $\mathcal{L}_{\text{LPA-real}}$  for the LPA model), and  $\lambda_{\text{consist}}$  is a hyperparameter that controls the strength of consistency regularization. The student parameters  $\theta_S^{(r)}$  are updated by minimizing  $\mathcal{L}_M$ .

## C. Teacher Update (EMA).

*Goal.* After updating the student parameters  $\theta_S^{(r)}$ , we update the teacher parameters  $\theta_T^{(r)}$  via an exponential moving average to obtain a smoother, more stable model that will be used in the next E-step.

For each parameter vector (e.g., for the fusion or LPA model) we apply:

$$\theta_T^{(r)} \leftarrow \alpha \cdot \theta_T^{(r-1)} + (1 - \alpha) \cdot \theta_S^{(r)},$$

where  $\alpha \in (0, 1)$  is a high momentum coefficient, typically close to 1 (e.g.,  $\alpha = 0.999$ ). This update ensures that the teacher evolves slowly and acts as a low-variance target for pseudo-label generation.

## D. Stopping Criteria.

The EM-refinement loop over  $r$  is terminated based on a combination of label stability and validation performance:

- *Label stability:* Let  $\mathcal{H}^{(r)}$  be the set of high-confidence links in  $C_{\text{prior}}^{(r)}$  (e.g., links with  $w_{ij}$  above a fixed threshold). We compute the Jaccard similarity

$$J^{(r)} = \frac{|\mathcal{H}^{(r)} \cap \mathcal{H}^{(r-1)}|}{|\mathcal{H}^{(r)} \cup \mathcal{H}^{(r-1)}|}$$

and stop if  $J^{(r)}$  exceeds a threshold (e.g., 0.99), indicating that the answer key has stabilized.

- *Validation plateau:* We monitor the total loss  $\mathcal{L}_M$  and other relevant metrics (e.g., calibration error) on a held-out synthetic validation set from Step 4.1. Training stops if these metrics do not improve for a fixed number of rounds.

### Final Output of EM Refinement

After convergence, the final teacher parameter sets  $\theta_T^{(\text{final})}$  define the refined models used in the online CoCaD inference pipeline. Concretely, we obtain:

$$f_{\text{fusion}}^{(\text{final})}, \quad LPA_{\text{model}}^{(\text{final})}, \quad CPC_{\text{model}}^{(\text{final})}, \quad f_{\text{student}}^{(\text{final})},$$

each of which has been initialized on synthetic data and subsequently adapted to the real corpus via the EM-refinement procedure described above.